aggregateByKey() Transformation

Syntax of this transformation

source_rdd.aggregateByKey(zeroValue, lambda1, lambda2) --> target_rdd

aggregateByKey(zeroValue, seqFunc, combFunc, numPartitions=None, partitionFunc=)

Aggregate the values of each key, using given combine functions and a neutral "zero value". This function can return a different result type, U, than the type of the values in this RDD, V. Thus, we need one operation for merging a V into a U and one operation for merging two U's. The former operation is used for merging values within a partition, and the latter is used for merging values between partitions. To avoid memory allocation, both of these functions are allowed to modify and return their first argument instead of creating a new U.

In [1]:
```python
from pyspark.sql import SparkSession

spark = SparkSession \
        .builder \
        .master("local[*]") \
        .appName("groupByKey Transformation") \
        .enableHiveSupport() \
        .getOrCreate()

#path of the data file on the local machine
data_file = '/Users/vaishaliyasala/Desktop/Github/Spark/Exercise_Dependencies/sales_data.csv'

#Read the csv into a dataframe
df = spark.read.csv(data_file, header = True, )

df1 = df.select(df["InvoiceNo"],df["UnitPrice"],df["Quantity"]).repartition(4)

print(df1.printSchema())

#Creating view of the dataframe of with 3 required columns and sample of 3% of data
sample_df = df1.sample(0.02,134)

sample_df.show()
```

```
22/10/13 16:47:57 WARN Utils: Your hostname, Vaishalis-MacBook-Pro.local resolves to a loopback address: 127.0.
0.1; using 192.168.0.105 instead (on interface en0)
22/10/13 16:47:57 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
22/10/13 16:47:57 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builti
n-java classes where applicable
```

```
root
 |-- InvoiceNo: string (nullable = true)
 |-- UnitPrice: string (nullable = true)
 |-- Quantity: string (nullable = true)

None
+---------+---------+--------+
|InvoiceNo|UnitPrice|Quantity|
+---------+---------+--------+
|   536464|     2.55|       1|
|   536408|     0.65|      36|
|   536412|     1.65|       5|
|   536412|     1.65|       3|
|   536464|     1.95|       1|
|   536415|     2.95|       3|
|   536399|     1.85|       6|
|   536401|     5.95|       1|
|   536409|     0.65|      12|
|   536520|      2.1|       2|
|   536409|     2.95|       1|
|   536392|      165|       1|
|   536414|        0|      56|
|   536464|     1.25|       3|
|   536420|     2.95|       6|
|   536396|     1.06|       6|
|   536520|     1.95|       5|
|   536389|     4.95|       8|
|   536446|     0.42|      10|
|   536375|     6.95|       4|
+---------+---------+--------+
only showing top 20 rows
```

In [2]:
```python
# apply a map() transformation to rdd to create (K, V) pairs

#In this key-value pair, key is the InvoiceNo and the number is the value

#whereas the price is obtained from UnitPrice*Qunatity


import decimal


def get_price(x3):
    try:
        UnitPrice = decimal.Decimal(x3[2])
        convert = UnitPrice * decimal.Decimal(x3[1])
    except decimal.InvalidOperation:
            print("Invalid input")
    key = x3[0]
    price = convert
    return (key, price)


new_rdd = df1.rdd.map(lambda x : get_price(x))
print("Number of elements =",len(new_rdd.collect()))
print("Number of Partitions =",new_rdd.getNumPartitions())
```

```
Number of elements = 999
Number of Partitions = 4
```

In [3]:
```python
#Showing the Result for the dataframe sample sample_df
sample_df_rdd = sample_df.rdd.map(lambda x : get_price(x))

print(sample_df_rdd.collect())
```

```
[('536464', Decimal('2.55')), ('536408', Decimal('23.40')), ('536412', Decimal('8.25')), ('536412', Decimal('4.
95')), ('536464', Decimal('1.95')), ('536415', Decimal('8.85')), ('536399', Decimal('11.10')), ('536401', Decim
al('5.95')), ('536409', Decimal('7.80')), ('536520', Decimal('4.2')), ('536409', Decimal('2.95')), ('536392', D
ecimal('165')), ('536414', Decimal('0')), ('536464', Decimal('3.75')), ('536420', Decimal('17.70')), ('536396',
Decimal('6.36')), ('536520', Decimal('9.75')), ('536389', Decimal('39.60')), ('536446', Decimal('4.20')), ('536
375', Decimal('27.80')), ('536373', Decimal('6.36')), ('536408', Decimal('9.90'))]]
```

In [4]:
```python
# rdd2 = rdd1.aggregateByKey(
#     zero_value,
#     lambda x,y: (x[0] + y,    x[1] + 1),
#     lambda x,y: (x[0] + y[0], x[1] + y[1])
# )

# Where the following is true about the meaning of each x and y
# pair above :
#
# First lambda expression for Within-Partition Reduction Step::
#     x: is a TUPLE that holds: (runningSum, runningCount).
#     y: is a SCALAR that holds the next Value

# Second lambda expression for Cross-Partition Reduction Step::
#     x: is a TUPLE that holds: (runningSum, runningCount).
#     y: is a TUPLE that holds: (nextPartitionsSum, nextPartitionsCount).


# we are showing for each key (invoice), U is sum of prices for all items and no of items
price_and_count = new_rdd.aggregateByKey(
        (0, 0), \
        lambda x, y: (x[0]+y, x[1]+1), \
        lambda rdd1, rdd2: (rdd1[0] + rdd2[0], rdd1[1] + rdd2[1]) \
    )


print("price_and_count.count() = ", price_and_count.count())
print("price_and_count.collect() = ", price_and_count.collect())
```

```
sum_count.count() =  66
sum_count.collect() =  [('536460', Decimal('295.54'), 14), ('C536391', (Decimal('-141.48'), 7)), ('536412',
(Decimal('514.41'), 81)), ('536401', (Decimal('354.23'), 64)), ('536420', (Decimal('233.45'), 14)), ('536385',
(Decimal('130.85'), 7)), ('536397', (Decimal('279.00'), 2)), ('536416', (Decimal('225.70'), 6)), ('536508', (De
cimal('155.52'), 2)), ('536376', (Decimal('328.80'), 2)), ('536405', (Decimal('326.40'), 1)), ('536393', (Decim
al('79.60'), 1)), ('536374', (Decimal('350.40'), 1)), ('536369', (Decimal('17.85'), 1)), ('536380', (Decimal('3
4.80'), 1)), ('536389', (Decimal('358.25'), 14)), ('536370', (Decimal('855.86'), 20)), ('536488', (Decimal('16
5.89'), 35)), ('536425', (Decimal('362.45'), 17)), ('536409', (Decimal('243.28'), 58)), ('536402', (Decimal('35
7.00'), 3)), ('536367', (Decimal('278.73'), 12)), ('536406', (Decimal('353.14'), 17)), ('536368', (Decimal('70.
05'), 4)), ('536365', (Decimal('139.12'), 7)), ('C536379', (Decimal('-27.5'), 1)), ('536502', (Decimal('95.2
9'), 5)), ('536371', (Decimal('204.00'), 1)), ('536414', (Decimal('0'), 1)), ('536400', (Decimal('17.40'), 1)),
('536375', (Decimal('259.86'), 16)), ('536390', (Decimal('1825.74'), 24)), ('536396', (Decimal('376.36'), 18)),
('536404', (Decimal('471.30'), 28)), ('536464', (Decimal('277.35'), 85)), ('536408', (Decimal('783.11'), 48)),
('536372', (Decimal('22.20'), 2)), ('536378', (Decimal('444.98'), 19)), ('536398', (Decimal('426.56'), 17)),
('536514', (Decimal('295.50'), 5)), ('536446', (Decimal('440.89'), 32)), ('536477', (Decimal('2474.74'), 14)),
('536382', (Decimal('430.60'), 12)), ('536463', (Decimal('17.40'), 1)), ('536377', (Decimal('22.20'), 2)), ('53
6387', (Decimal('3193.92'), 5)), ('536407', (Decimal('22.20'), 2)), ('536403', (Decimal('192.60'), 2)), ('536638
6', (Decimal('508.20'), 3)), ('C536383', (Decimal('-4.65'), 1)), ('536395', (Decimal('507.88'), 14)), ('536642
3', (Decimal('303.90'), 9)), ('536415', (Decimal('390.79'), 59)), ('536500', (Decimal('223.90'), 15)), ('536637
3', (Decimal('259.86'), 16)), ('536384', (Decimal('489.60'), 13)), ('536520', (Decimal('265.87'), 54)), ('536638
1', (Decimal('449.98'), 35)), ('536388', (Decimal('226.14'), 14)), ('536392', (Decimal('318.14'), 10)), ('536639
4', (Decimal('1024.68'), 11)), ('536399', (Decimal('22.20'), 2)), ('536366', (Decimal('22.20'), 2)), ('536437',
(Decimal('842.12'), 6)), ('536466', (Decimal('42.90'), 2)), ('C536506', (Decimal('-25.50'), 1))]
```