

coalesce() Transformation

In [1]:

```
from pyspark.sql import SparkSession
import pyspark

spark = SparkSession \
    .builder \
    .master("local[*]") \
    .appName("coalesce Transformation") \
    .enableHiveSupport() \
    .getOrCreate()
```

22/10/14 12:05:42 WARN Utils: Your hostname, Vaishalis-MacBook-Pro.local resolves to a loopback address: 127.0.0.1; using 192.168.0.105 instead (on interface en0)

22/10/14 12:05:42 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address

Setting default log level to "WARN".

To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).

22/10/14 12:05:43 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

22/10/14 12:05:44 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 4041.

22/10/14 12:05:44 WARN Utils: Service 'SparkUI' could not bind on port 4041. Attempting port 4042.

In [2]:

```
def debug_a_partition(iterator):
    print("==begin-partition=")
    for x in iterator:
        print(x)
    #end-for
    print("==end-partition=")

names_list = ["ABC - 1", "DEF - 1","GHI-1", "ABC - 2", "DEF - 2","GHI -2", \
              "ABC - 3", "DEF - 3","GHI -3","ABC - 4", "DEF - 4","GHI -4"]

names_rdd = spark.sparkContext.parallelize(names_list,4)

print("From local[4] =",names_rdd.getNumPartitions())
print("Repartition elements : ", names_rdd.foreachPartition(debug_a_partition))
```

From local[4] = 4

[Stage 0:> (0 + 4) / 4]

Repartition elements : None

==begin-partition==begin-partition=

==begin-partition=

==begin-partition=ABC - 2

DEF - 2ABC - 3

ABC - 4

DEF - 3

DEF - 4GHI -3

GHI -4GHI -2

==end-partition=

==end-partition===end-partition=

ABC - 1

DEF - 1

GHI-1

==end-partition=

In [3]:

```
#repartition shuffles the data completely
repartition_rdd = names_rdd.repartition(3)
print("coalesce elements : ", repartition_rdd.foreachPartition(debug_a_partition))
```

[Stage 1:> (0 + 4) / 4]

coalesce elements : None

==begin-partition=

==end-partition=

==begin-partition=

ABC - 2

DEF - 2

GHI -2

ABC - 4

DEF - 4

GHI -4

==end-partition=

==begin-partition=

ABC - 1

DEF - 1

GHI-1

ABC - 3

DEF - 3

GHI -3

==end-partition=

From the output above, we can see the reshuffling of the entire data from 4 partitions in names_rdd to 3 partitions in repartition_rdd. also, partition 1 is empty.

In [4]:

```
#coalesce combines the partition close to it
coalesce_rdd = names_rdd.coalesce(3)
print("Repartition size : ", coalesce_rdd.foreachPartition(debug_a_partition))
```

==begin-partition=

==begin-partition=ABC - 1

DEF - 1

GHI-1

==begin-partition===end-partition=ABC - 2

DEF - 2

GHI -2

==end-partition=

ABC - 3

DEF - 3

GHI -3

ABC - 4

DEF - 4

GHI -4

==end-partition=

Repartition size : None

From the output above, we can see partitions 3 and 4 of names_rdd combined to form partition 3 in coalesce_rdd rather than shuffle around the entire data.

In [5]:

```
#Reading data from a file on the local machine
data_file_path = "/Users/vaishaliyasala/Desktop/Github/Spark/Exercise_Dependencies/SalesJan2009.csv"

df = spark.read.csv(data_file_path, header = True )

df1 = df.select(df["Name"],df["Country"]).repartition(4)

print("Count in the original data=", df1.count())

#Filter the names only from country United States
filtered_rdd = df1.rdd.filter(lambda x: (x[1] == "United States"))

filtered_rdd.toDF(["Name", "Country"]).show(5)

print("Filtered data Count =", filtered_rdd.count())
print("Number of Partitions =", filtered_rdd.getNumPartitions())
```

Count in the original data= 998

+-----+-----+

| Name| Country|

+-----+-----+

| Abikay|United States|

|Christian|United States|

| Alicja |United States|

| Debora |United States|

| Sandrine|United States|

+-----+-----+

only showing top 5 rows

Filtered data Count = 463

Number of Partitions = 4

In [6]:

```
#As the filtered data count increased significantly, we can reduce the number of partitions from 4 to 2
#It doesn't change the resultant RDD as seen from the outputs of both before and after coalesce transformation
#Only data changes in each partition

names_coalesce_rdd = filtered_rdd.coalesce(2)

names_coalesce_rdd.toDF(["Name", "Country"]).show(5)

print("Number of Partitions =", names_coalesce_rdd.getNumPartitions())
```

+-----+-----+

| Name| Country|

+-----+-----+

| Abikay|United States|

|Christian|United States|

| Alicja |United States|

| Debora |United States|

| Sandrine|United States|

+-----+-----+

only showing top 5 rows

Number of Partitions = 2