

reduceByKey() Transformation

groupByKey() Transformation cannot be used on large datasets. To solve this problem, let us look us at reduceByKey() Transformation. reduceByKey is optimized with a map side combine. This means it performs the merging locally on each mapper for each key before sending results to a reducer operation. After that, the values are combined for each key using an associative and commutative reduce function. By this, less elements are sent over the network.

In [1]:

```
from pyspark.sql import SparkSession
import pyspark

spark = SparkSession \
    .builder \
    .master("local[4]") \
    .appName("Sample Transformation") \
    .enableHiveSupport() \
    .getOrCreate()

#path of the data file on the local machine
data_file = '/Users/vaishaliyasala/Desktop/Github/Spark/Exercise_Dependencies/sales_data.csv'

#Read the csv into a dataframe
df = spark.read.csv(data_file, header = True )

df1 = df.select(df["InvoiceNo"],df["UnitPrice"],df["Quantity"]).repartition(4)

print(df1.printSchema())

#Creating view of the dataframe of with 3 required columns and sample of 2% of data
sample_df = df1.sample(0.02,134)

sample_df.show()
```

22/10/13 16:40:49 WARN Utils: Your hostname, Vaishalis-MacBook-Pro.local resolves to a loopback address: 127.0.0.1; using 192.168.0.105 instead (on interface en0)
22/10/13 16:40:49 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
22/10/13 16:40:50 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using built-in-java classes where applicable
22/10/13 16:40:51 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 4041.

```
root
 |-- InvoiceNo: string (nullable = true)
 |-- UnitPrice: string (nullable = true)
 |-- Quantity: string (nullable = true)

None
```

InvoiceNo	UnitPrice	Quantity
536464	2.55	1
536408	0.65	36
536412	1.65	5
536412	1.65	3
536464	1.95	1
536415	2.95	3
536399	1.85	6
536401	5.95	1
536409	0.65	12
536520	2.1	2
536409	2.95	1
536392	165	1
536414	0	56
536464	1.25	3
536420	2.95	6
536396	1.06	6
536520	1.95	5
536389	4.95	8
536446	0.42	10
536375	6.95	4

only showing top 20 rows

In [2]:

```
# apply a map() transformation to rdd to create (K, V) pairs

#In this key-value pair, key is the InvoiceNo and the number is the value

#whereas the price is obtained from UnitPrice*Qunatity

import decimal

def get_price(x3):
    try:
        UnitPrice = decimal.Decimal(x3[2])
        convert = UnitPrice * decimal.Decimal(x3[1])
    except decimal.InvalidOperation:
        print("Invalid input")
    key = x3[0]
    price = convert
    return (key, price)

rdd1 = df1.rdd.map(lambda x : get_price(x))
print("Number of elements =",len(rdd1.collect()))
print("Number of Partitions =",rdd1.getNumPartitions())

#Showing the Result for the dataframe sample sample_df
sample_df_rdd = sample_df.rdd.map(lambda x : get_price(x))

print(sample_df_rdd.collect())
```

Number of elements = 999
Number of Partitions = 4
[('536464', Decimal('2.55')), ('536408', Decimal('23.40')), ('536412', Decimal('8.25')), ('536412', Decimal('4.95')), ('536464', Decimal('1.95')), ('536415', Decimal('8.85')), ('536399', Decimal('11.10')), ('536401', Decimal('5.95')), ('536409', Decimal('7.80')), ('536520', Decimal('4.2')), ('536409', Decimal('2.95')), ('536392', Decimal('165')), ('536414', Decimal('0')), ('536464', Decimal('3.75')), ('536420', Decimal('17.70')), ('536396', Decimal('6.36')), ('536520', Decimal('9.75')), ('536389', Decimal('39.60')), ('536446', Decimal('4.20')), ('536375', Decimal('27.80')), ('536373', Decimal('6.36')), ('536408', Decimal('9.90'))]

In [3]:

```
# apply a reduceByKey() transformation on rdd1 to create a (key, value) pair

# where key is the InvoiceNo and value is sum of prices for each key

#we can create more partitions than its parent RDD.

rdd2 = rdd1.reduceByKey(lambda a, b: (a+b),10)

print(len(rdd2.collect()),rdd2.getNumPartitions())
print("Number of elements =",len(rdd2.collect()))
print("Number of Partitions =",rdd2.getNumPartitions())
```

66 10
Number of elements = 66
Number of Partitions = 10

From result of Input block 2 and 3, we can see the number of elements decreased because they are merged together when they have the same key. Additionally, it is optimized with a map side combine.

In [4]:

```
#Below we can the result of reduceByKey() applied on sample_df_rdd

print(sample_df_rdd.reduceByKey(lambda a, b: (a+b),10).collect())
```

[('536464', Decimal('8.25')), ('536396', Decimal('6.36')), ('536399', Decimal('11.10')), ('536520', Decimal('13.95')), ('536375', Decimal('27.80')), ('536414', Decimal('0')), ('536373', Decimal('6.36')), ('536389', Decimal('39.60')), ('536412', Decimal('13.20')), ('536409', Decimal('10.75')), ('536408', Decimal('33.30')), ('536401', Decimal('5.95')), ('536420', Decimal('17.70')), ('536446', Decimal('4.20')), ('536415', Decimal('8.85')), ('536392', Decimal('165'))]

In []: